

BROCODE SOLUTIONS

AI SOFTWARE DEVELOPMENT

Abu Dhabi · United Arab Emirates

A Brocode AI Insights publication

GulfBench v2 — LLM Fine-Tuning Evaluation Pack

An open eval suite for Khaleeji-tuned LLMs with reproducer code.

Inside this report

- 01 Benchmark methodology
- 02 Base model comparison (Mistral, Llama, Jais)
- 03 Fine-tuning recipes (QLoRA, DPO, ORPO)
- 04 Cost and latency at H100 / H200
- 05 Reproducer on GitHub

brocode.ae · hello@brocode.ae · +971 50 761 2213

Benchmark methodology

This is a preview of the "Benchmark methodology" section of the report. The final published version of this report contains in-depth analysis, benchmark data, reference architectures, and recommendations specific to enterprises operating in the UAE and the wider GCC. If you would like the full version before it is released, please request a copy from the Brocode AI team and we will share an embargoed draft under NDA. We can also arrange a 30-minute briefing with the engineer who authored this section. This preview is intentionally short so that the reader of the gated download can scan the structure of the full report. The body sections of the final publication run between 6 and 12 pages each, with a mix of prose, tables, and figures.

Base model comparison (Mistral, Llama, Jais)

This is a preview of the "Base model comparison (Mistral, Llama, Jais)" section of the report. The final published version of this report contains in-depth analysis, benchmark data, reference architectures, and recommendations specific to enterprises operating in the UAE and the wider GCC. If you would like the full version before it is released, please request a copy from the Brocode AI team and we will share an embargoed draft under NDA. We can also arrange a 30-minute briefing with the engineer who authored this section. This preview is intentionally short so that the reader of the gated download can scan the structure of the full report. The body sections of the final publication run between 6 and 12 pages each, with a mix of prose, tables, and figures.

Fine-tuning recipes (QLoRA, DPO, ORPO)

This is a preview of the "Fine-tuning recipes (QLoRA, DPO, ORPO)" section of the report. The final published version of this report contains in-depth analysis, benchmark data, reference architectures, and recommendations specific to enterprises operating in the UAE and the wider GCC. If you would like the full version before it is released, please request a copy from the Brocode AI team and we will share an embargoed draft under NDA. We can also arrange a 30-minute briefing with the engineer who authored this section. This preview is intentionally short so that the reader of the gated download can scan the structure of the full report. The body sections of the final publication run between 6 and 12 pages each, with a mix of prose, tables, and figures.

Cost and latency at H100 / H200

This is a preview of the "Cost and latency at H100 / H200" section of the report. The final published version of this report contains in-depth analysis, benchmark data, reference architectures, and recommendations specific to enterprises operating in the UAE and the wider GCC. If you would like the full version before it is released, please request a copy from the Brocode AI team and we will share an embargoed draft under NDA. We can also arrange a 30-minute briefing with the engineer who authored this section. This preview is intentionally short so that the reader of the gated download can scan the structure of the full report. The body sections of the final publication run between 6 and 12 pages each, with a mix of prose, tables, and figures.

Reproducer on GitHub

This is a preview of the "Reproducer on GitHub" section of the report. The final published version of this report contains in-depth analysis, benchmark data, reference architectures, and recommendations specific to enterprises operating in the UAE and the wider GCC. If you would like the full version before it is released, please request a copy from the Brocode AI team and we will share an embargoed draft under NDA. We can also arrange a 30-minute briefing with the engineer who authored this section. This preview is intentionally short so that the reader of the gated download can scan the structure of the full report. The body sections of the final publication run between 6 and 12 pages each, with a mix of prose, tables, and figures.

Talk to our team

If this report sparked questions, we are happy to walk you through the underlying analysis, share the full benchmark dataset under NDA, or design a tailored architecture review for your specific environment. Reach us in whichever channel works.

hello@brocode.ae

+971 50 761 2213 (WhatsApp)

Al Maryah Island, Abu Dhabi Global Market, UAE

<https://www.brocode.ae>