

BROCODE SOLUTIONS

AI SOFTWARE DEVELOPMENT

Abu Dhabi · United Arab Emirates

A Brocode AI Insights publication

Sovereign LLM 36-month TCO Pack

Hardware BoM, power, cooling, and licence for a customer-owned LLM appliance.

Inside this report

- 01 Hardware BoM (8× H100 SXM5)
- 02 Power and cooling sizing
- 03 vLLM + Triton + TensorRT-LLM stack
- 04 36-month TCO vs Azure OpenAI vs OpenAI Enterprise
- 05 Customer Thales / Entrust HSM patterns

brocode.ae · hello@brocode.ae · +971 50 761 2213

Hardware BoM (8× H100 SXM5)

This is a preview of the "Hardware BoM (8× H100 SXM5)" section of the report. The final published version of this report contains in-depth analysis, benchmark data, reference architectures, and recommendations specific to enterprises operating in the UAE and the wider GCC. If you would like the full version before it is released, please request a copy from the Brocode AI team and we will share an embargoed draft under NDA. We can also arrange a 30-minute briefing with the engineer who authored this section. This preview is intentionally short so that the reader of the gated download can scan the structure of the full report. The body sections of the final publication run between 6 and 12 pages each, with a mix of prose, tables, and figures.

Power and cooling sizing

This is a preview of the "Power and cooling sizing" section of the report. The final published version of this report contains in-depth analysis, benchmark data, reference architectures, and recommendations specific to enterprises operating in the UAE and the wider GCC. If you would like the full version before it is released, please request a copy from the Brocode AI team and we will share an embargoed draft under NDA. We can also arrange a 30-minute briefing with the engineer who authored this section. This preview is intentionally short so that the reader of the gated download can scan the structure of the full report. The body sections of the final publication run between 6 and 12 pages each, with a mix of prose, tables, and figures.

vLLM + Triton + TensorRT-LLM stack

This is a preview of the "vLLM + Triton + TensorRT-LLM stack" section of the report. The final published version of this report contains in-depth analysis, benchmark data, reference architectures, and recommendations specific to enterprises operating in the UAE and the wider GCC. If you would like the full version before it is released, please request a copy from the Brocode AI team and we will share an embargoed draft under NDA. We can also arrange a 30-minute briefing with the engineer who authored this section. This preview is intentionally short so that the reader of the gated download can scan the structure of the full report. The body sections of the final publication run between 6 and 12 pages each, with a mix of prose, tables, and figures.

36-month TCO vs Azure OpenAI vs OpenAI Enterprise

This is a preview of the "36-month TCO vs Azure OpenAI vs OpenAI Enterprise" section of the report. The final published version of this report contains in-depth analysis, benchmark data, reference architectures, and recommendations specific to enterprises operating in the UAE and the wider GCC. If you would like the full version before it is released, please request a copy from the Brocode AI team and we will share an embargoed draft under NDA. We can also arrange a 30-minute briefing with the engineer who authored this section. This preview is intentionally short so that the reader of the gated download can scan the structure of the full report. The body sections of the final publication run between 6 and 12 pages each, with a mix of prose, tables, and figures.

Customer Thales / Entrust HSM patterns

This is a preview of the "Customer Thales / Entrust HSM patterns" section of the report. The final published version of this report contains in-depth analysis, benchmark data, reference architectures, and recommendations specific to enterprises operating in the UAE and the wider GCC. If you would like the full version before it is released, please request a copy from the Brocode AI team and we will share an embargoed draft under NDA. We can also arrange a 30-minute briefing with the engineer who authored this section. This preview is intentionally short so that the reader of the gated download can scan the structure of the full report. The body sections of the final publication run between 6 and 12 pages each, with a mix of prose, tables, and figures.

Talk to our team

If this report sparked questions, we are happy to walk you through the underlying analysis, share the full benchmark dataset under NDA, or design a tailored architecture review for your specific environment. Reach us in whichever channel works.

hello@brocode.ae

+971 50 761 2213 (WhatsApp)

Al Maryah Island, Abu Dhabi Global Market, UAE

<https://www.brocode.ae>